

OpenStack Day SP 2018

IPv6 no OpenStack

JOSÉ CASTILLO LEMA

DIRETORIA DE OPERAÇÕES E SOLUÇÕES AO CLIENTE



ESPECIALISTAS EM CLOUDS.

AGENDA

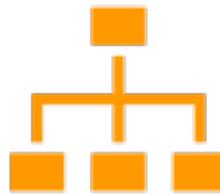
OPENSTACK DAY SP 2018 - IPV6 NO OPENSTACK



APRESENTAÇÃO
SOLUÇÃO
OPENSTACK



CRIAÇÃO DE
SUBREDES IPv6



CONEXÃO COM A
INTERNET



IMPLEMENTAÇÃO

MANDIC IN A NUTSHELL

Provedor público de nuvem

Modelo de serviço IaaS

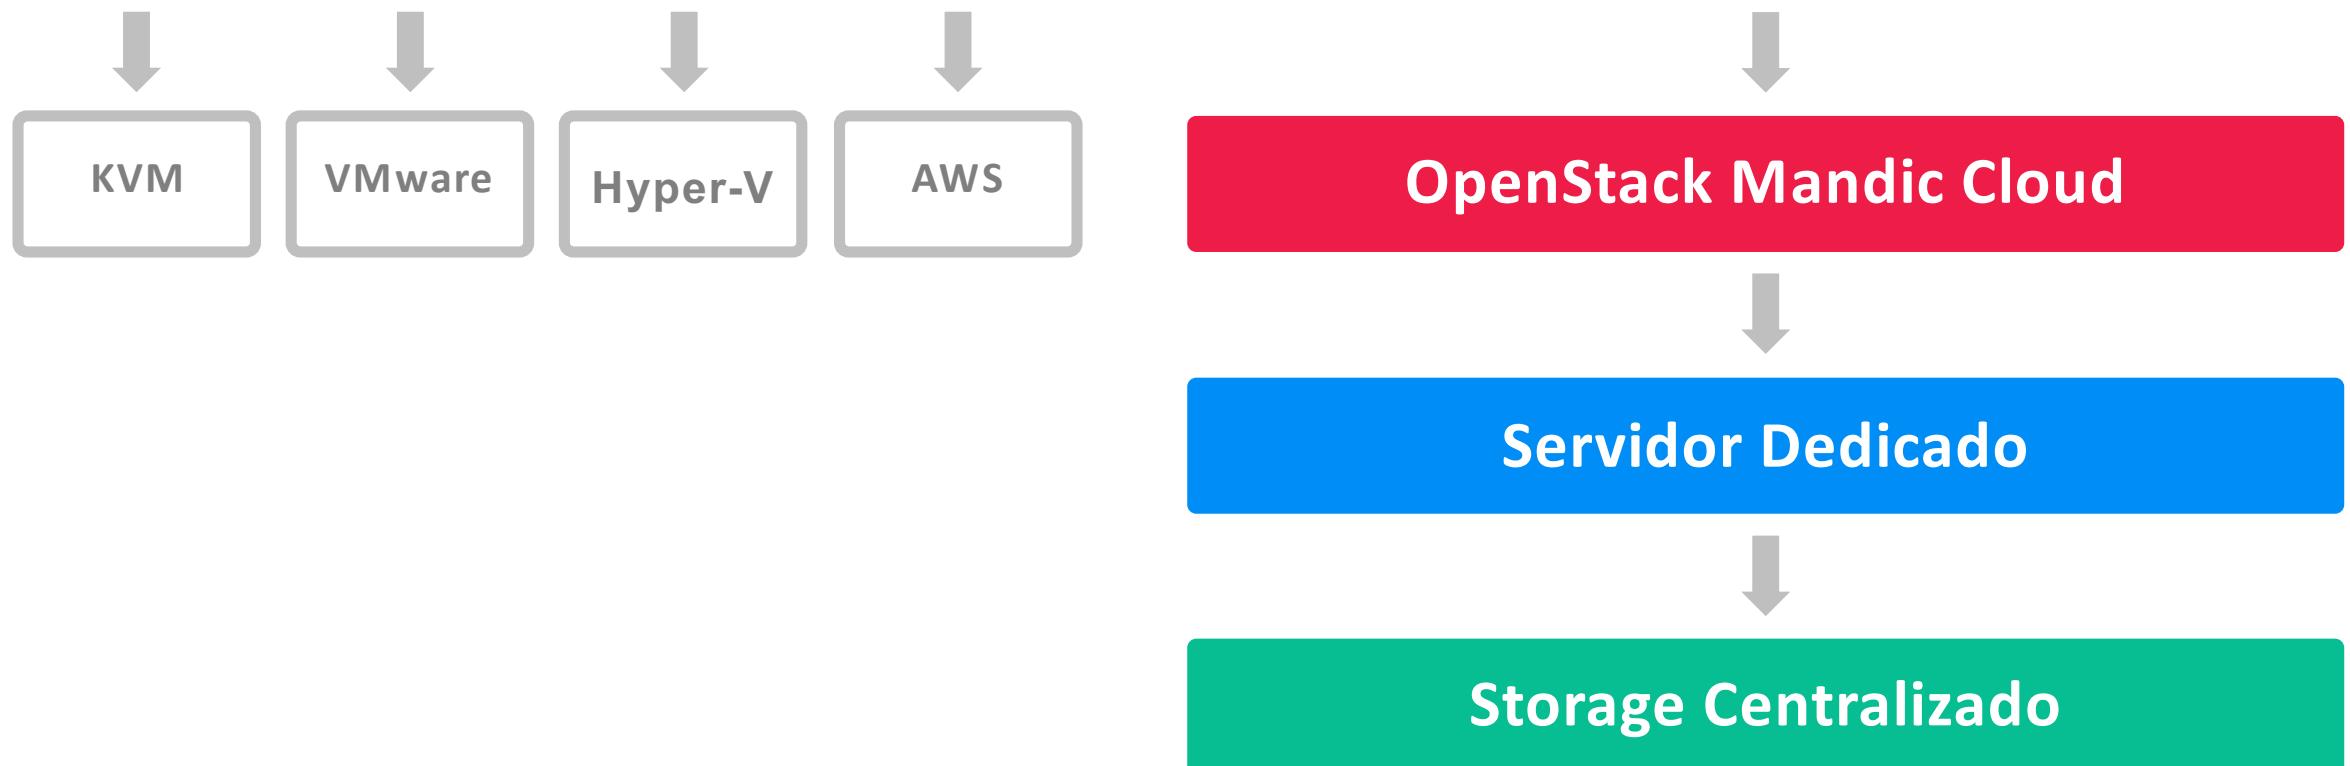
Forte foco empresarial



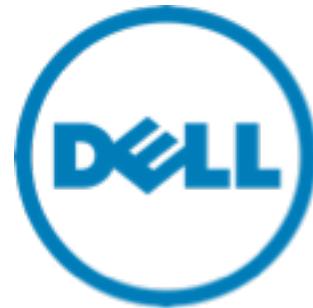
ESPECIALISTAS EM CLOUDS.

ARQUITETURA

PAINEL MANDIC CLOUD



PARCEIROS



redhat.

EQUIPE



GERD
JAKOBOVITSCH



SERGIO
DOS SANTOS OLIVEIRA



JOSE
CASTILLO LEMA

POR QUE IPv6?

- + **Abundância** de endereços IPv6
- + Elimina necessidade de *floating ips*
- + **Integração** de redes IPv6 de clientes
- + Elimina necessidade de NAT
- + Sem tráfego *broadcast*
- + Internet das Coisas

ONDE IMPLEMENTAR IPv6?

- + Servidores físicos
- + API – *endpoints*
- + Máquinas virtuais – *tenant networks*
 ⇒ *provider networks*

REQUISITOS DE IMPLEMENTAÇÃO

Ambiente ***dual-stack***. Todos os projetos/clientes com:

- + rede IPv4 /24 privada
- + rede IPv6 /64 pública
- + roteador virtual privado
 - + HA
 - + Sem DVR

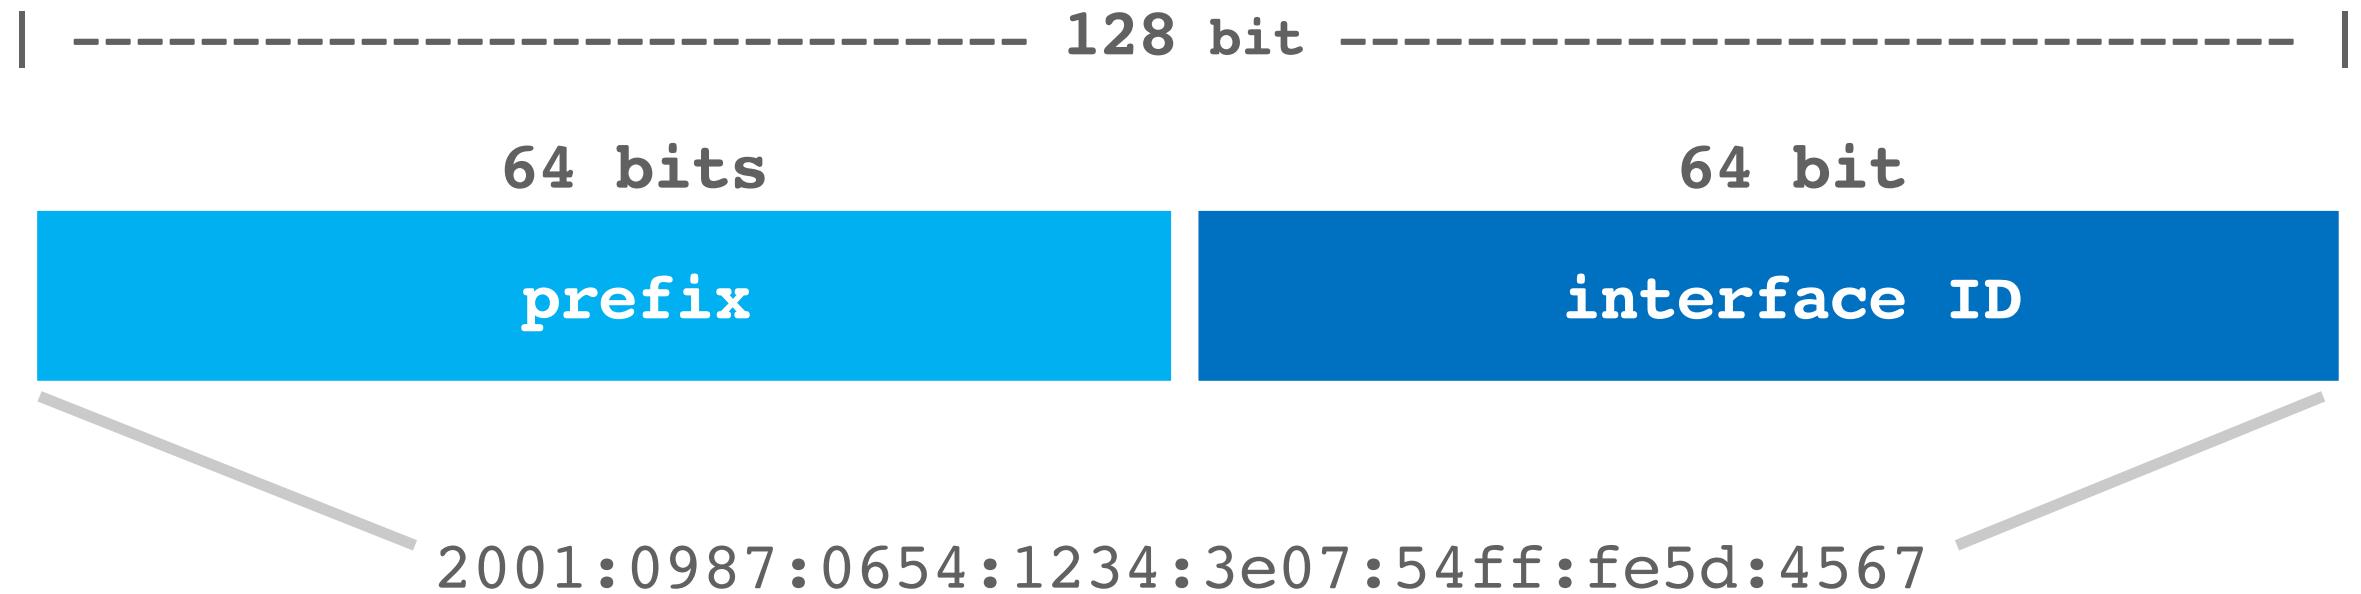
IPv6 NO OPENSTACK

CRIAÇÃO DE SUBREDES IPv6



IPv6 – ATRIBUIÇÃO DE ENDEREÇOS

SLAAC - Stateless Address Auto Configuration



IPv6 – ATRIBUIÇÃO DE ENDEREÇOS

DHCPv6 *stateless*

- + DNS
- + rotas

DHCPv6 *statefull*

IPv6 – ATRIBUIÇÃO DE ENDEREÇOS

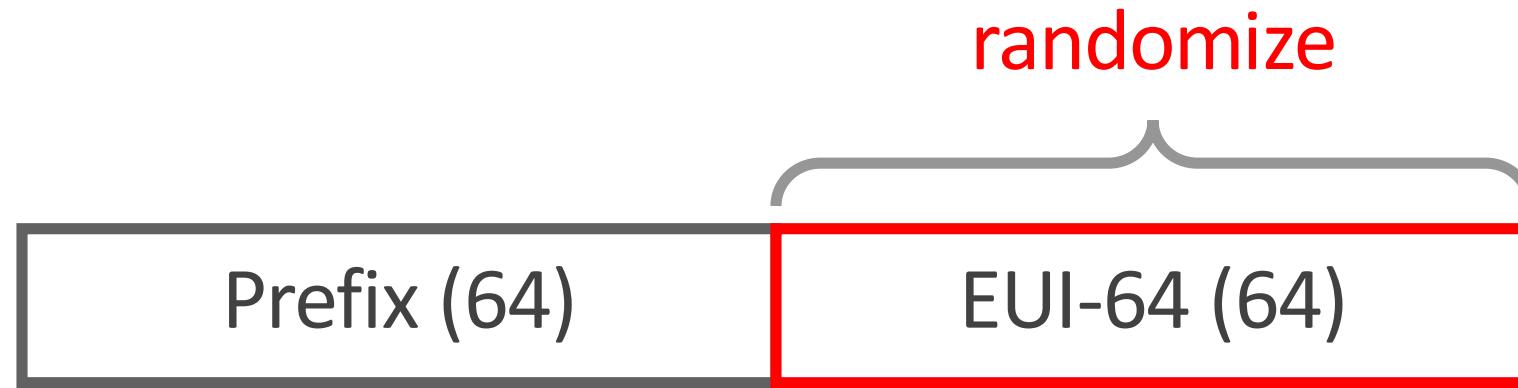
```
jose@jose-think:~$ openstack subnet show sub6-int
+-----+-----+
| Field      | Value
+-----+-----+
| allocation_pools | 2804:5d4:1:1000::2-2804:5d4:1:1000:ffff:ffff:ffff:ffff
| cidr        | 2804:5d4:1:1000::/64
| created_at   | 2018-03-12T18:30:55Z
| description    |
| dns_nameservers | 2001:4860:4860::8888
| enable_dhcp    | True
| gateway_ip     | 2804:5d4:1:1000::1
| host_routes    |
| id            | 2b1bffe8-0624-44f3-8f43-140d4fd1736d
| ip_version     | 6
| ipv6_address_mode | dhcpv6-stateless
| ipv6_ra_mode    | dhcpv6-stateless
| name          | sub6-int
| network_id     | 87130a10-be1a-493f-8f0f-511575ce7617
| project_id      | ed5f5ffe834244479ef3bc2bbc905fad
| revision_number | 2
| segment_id      | None
| service_types    |
| subnetpool_id   | 77f5924a-5058-4a5f-9a02-91d04ab131ce
| tags           |
| updated_at      | 2018-03-12T18:30:55Z
+-----+-----+
```

IPv6 – ATRIBUIÇÃO DE ENDEREÇOS

ip_version	6
ipv6_address_mode	dhcpv6-stateless
ipv6_ra_mode	dhcpv6-stateless

SLAAC: WINDOWS

- + Privacy Extensions for IPv6 SLAC



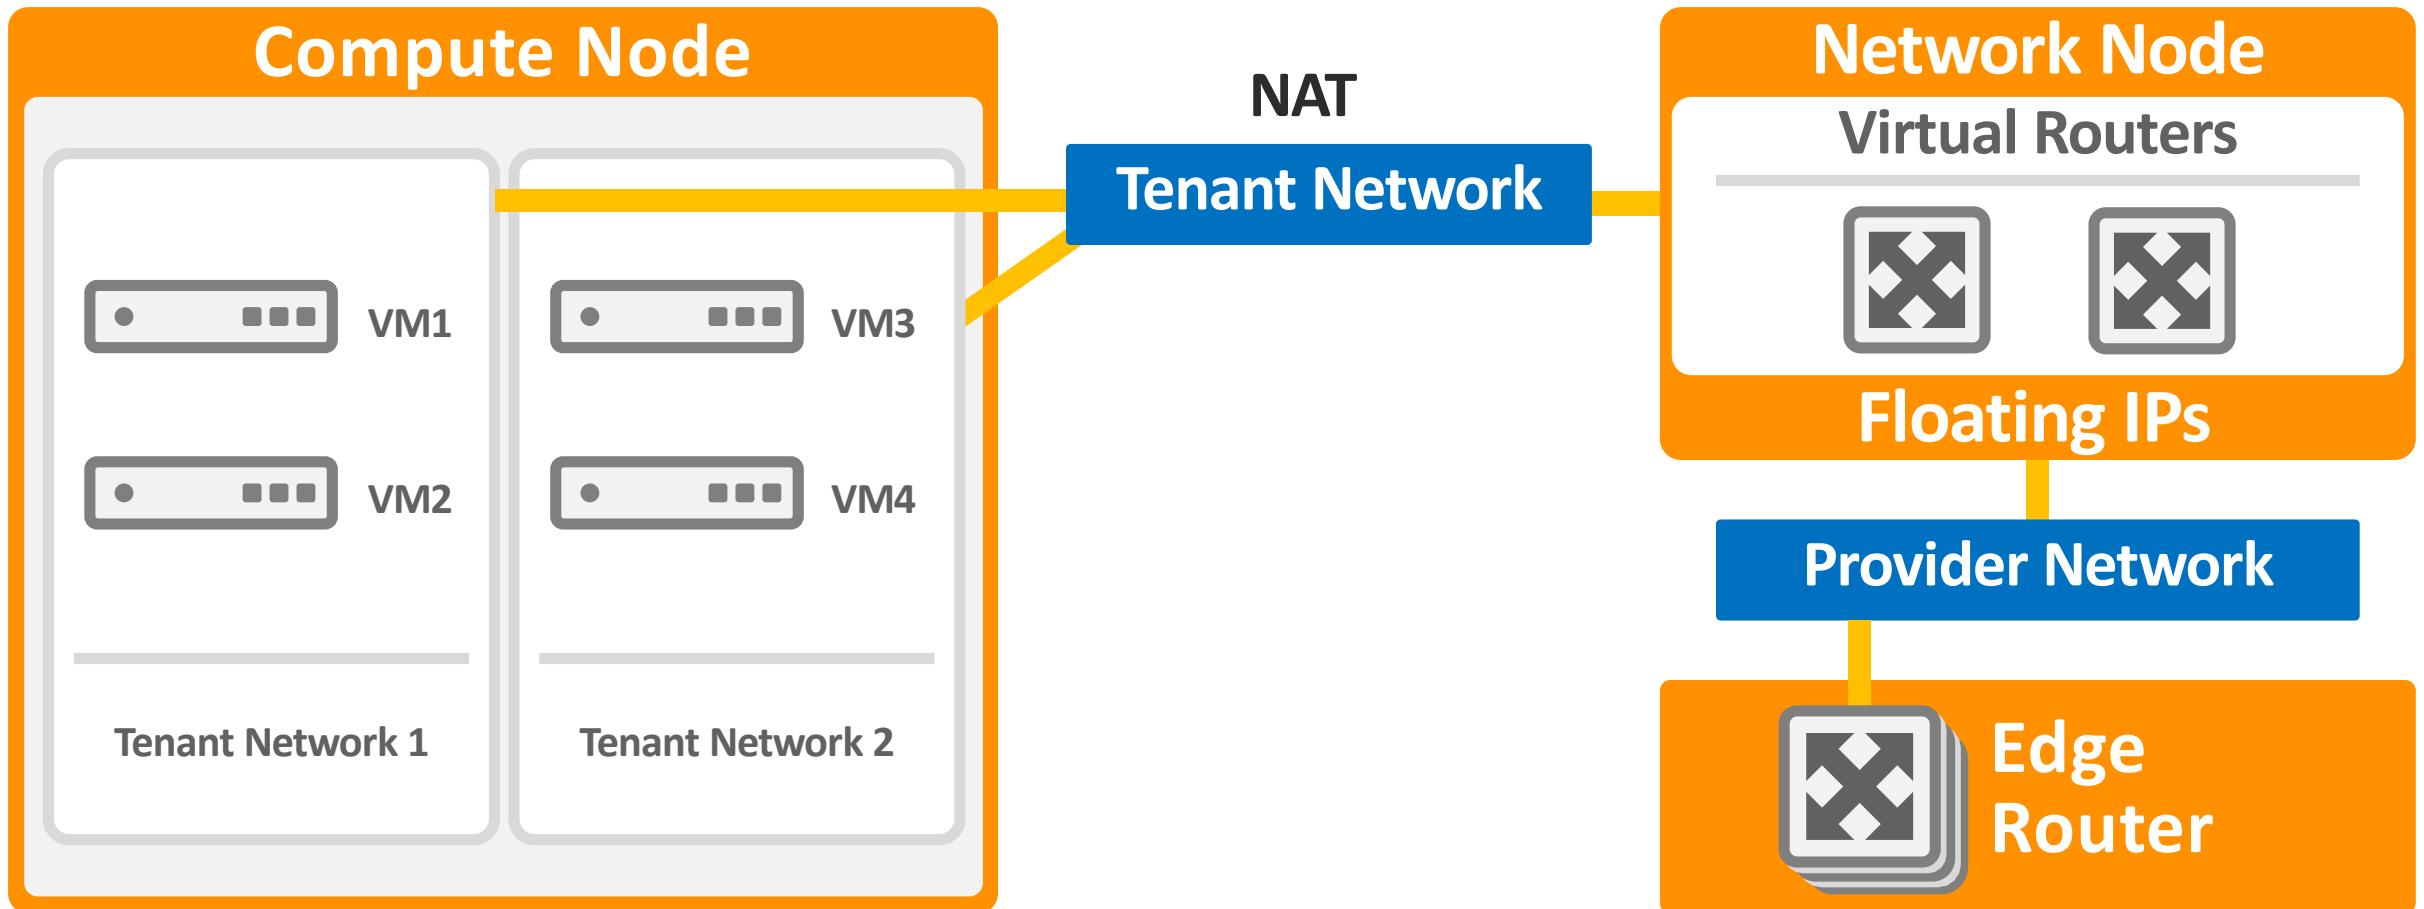
netsh interface ipv6 set global randomizeidentifiers=disabled

IPv6 NO OPENSTACK

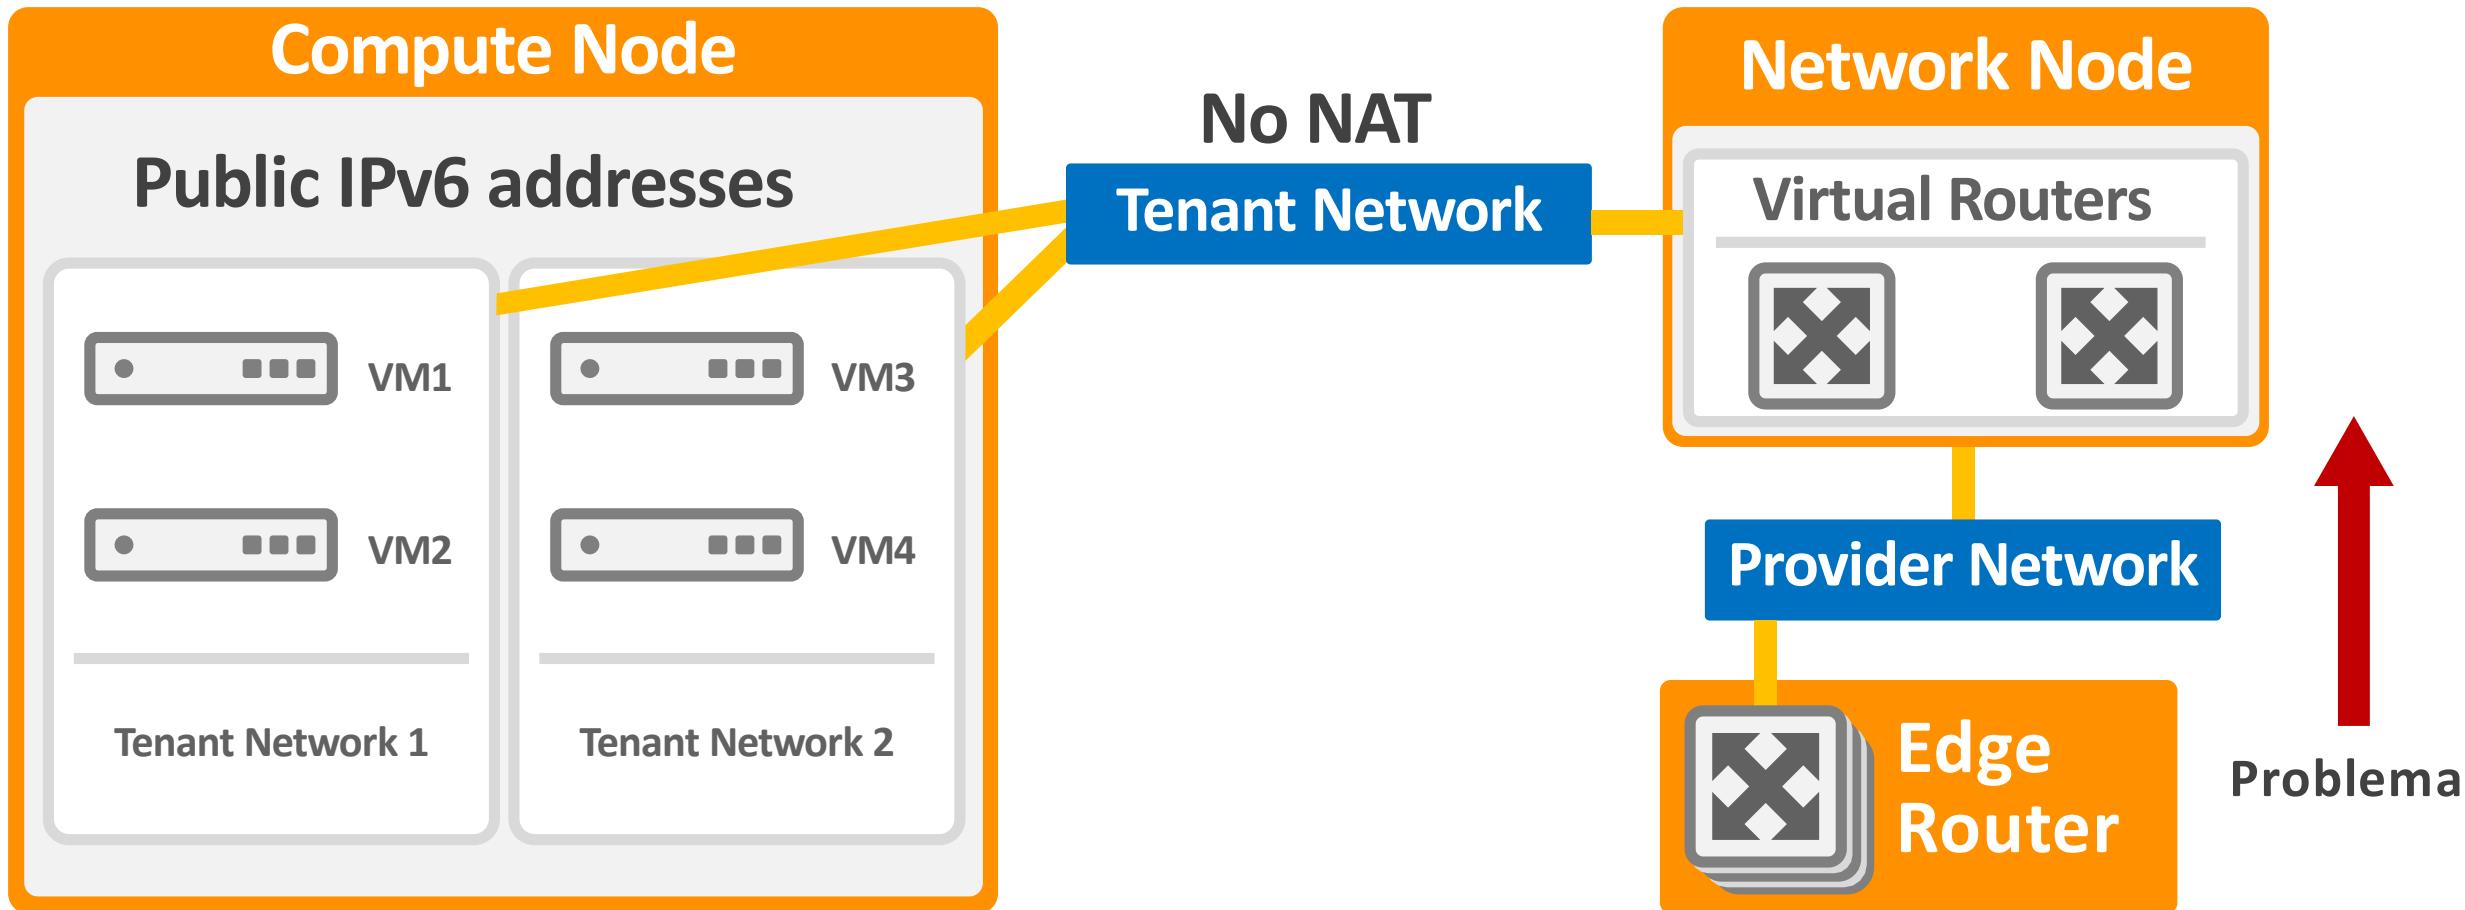
CONEXÃO COM A INTERNET



ARQUITETURA IPv4



ARQUITETURA IPv6



ARQUITETURA IPv6

Diferenças:

- + sem NAT
- + os endereços IPv6 públicos estão nas vms (nos compute nodes)
- + sem *floating IPs*, endereços IPv6 públicos são tratados como *fixed IPs* dentro do OpenStack

Desafio:

- + as vms conseguem sair para a Internet via IPv6
- + o roteador *edge* não sabe para qual roteador virtual encaminhar os pacotes

MODELOS DE IMPLEMENTAÇÃO

E-BGP usando **neutron-bgp-agent**

- + um agente BGP rodando nos *control nodes* atua como BGP *speaker*, propagando as rotas para o roteador *edge*
- + o roteador *edge* propaga as rotas para a Internet

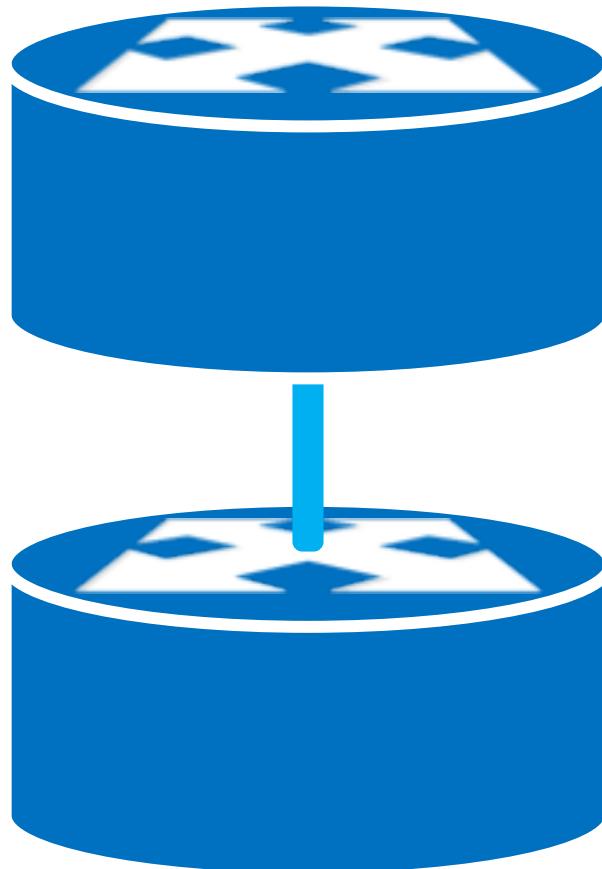
IPv6 *prefix delegation* usando **neutron-pd-agent**

- + o roteador *edge* atua como servidor DHCPv6
- + o roteador virtual solicita via DHCPv6 um prefixo de rede ao roteador *edge*

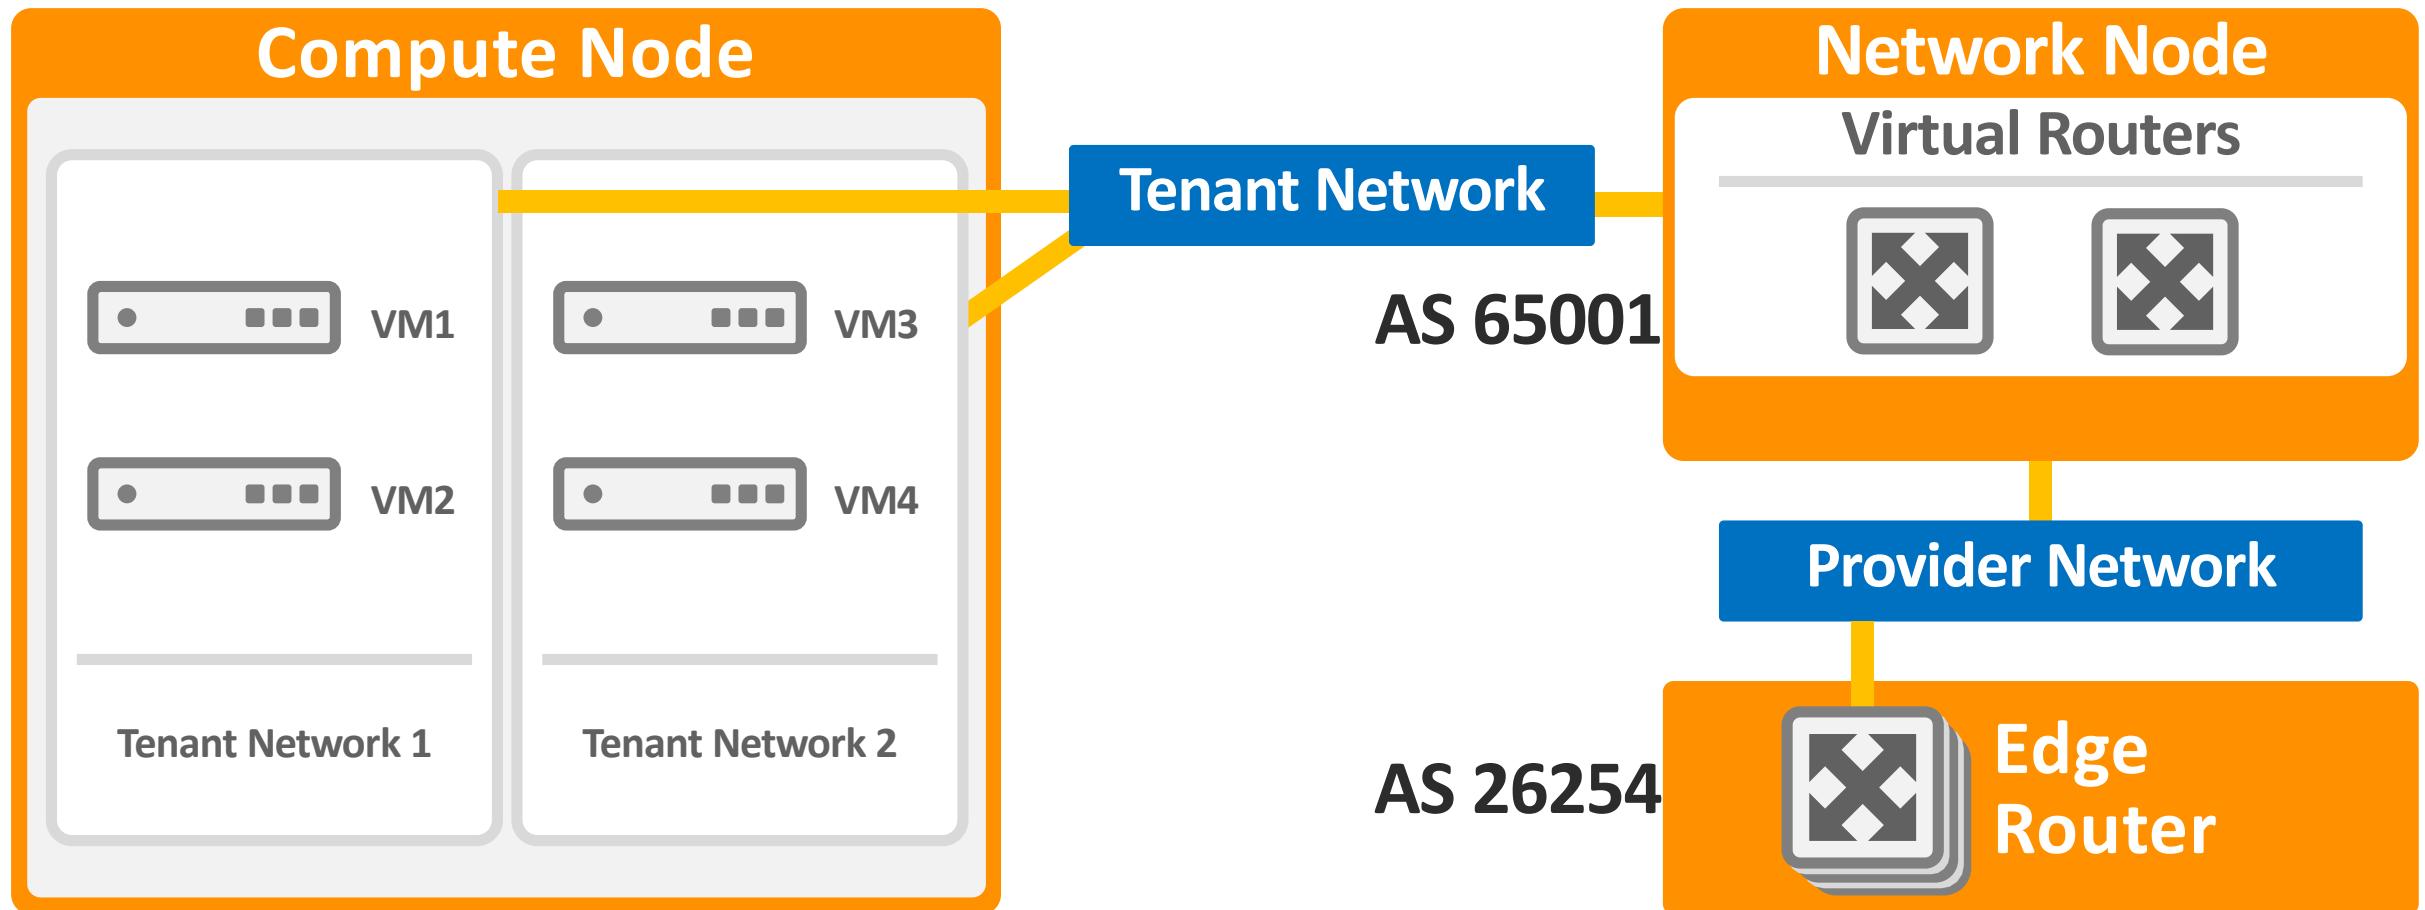
IPv6 PREFIX DELEGATION

Edge
router

Virtual
router



ROTEAMENTO



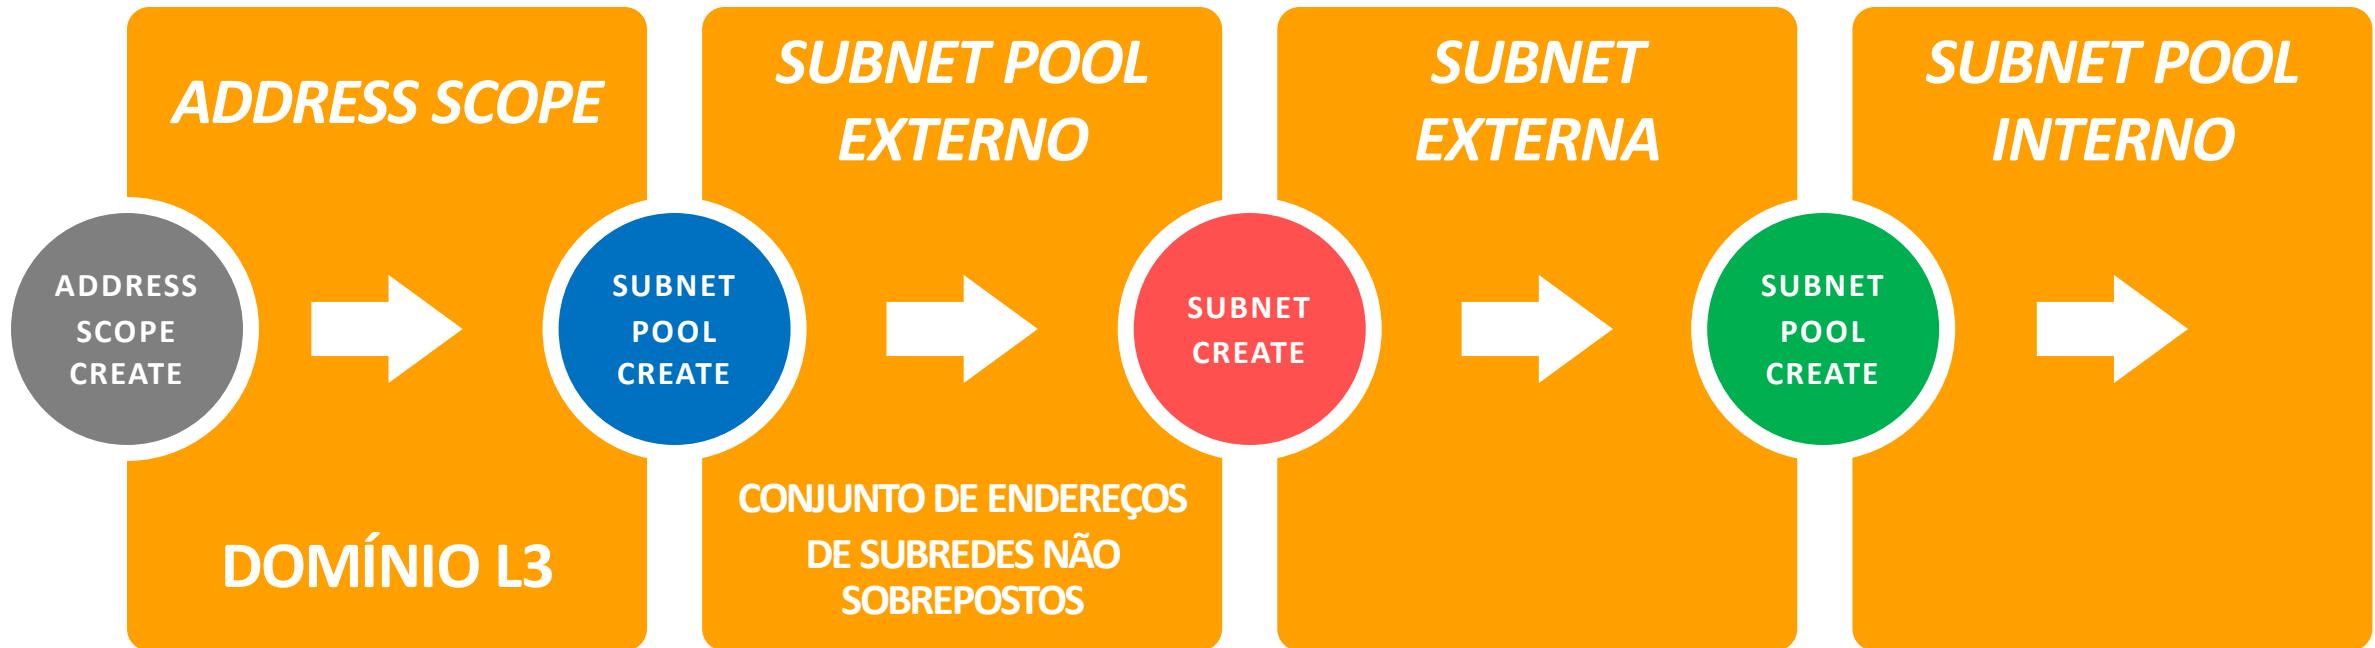
NEUTRON-BGP-AGENT

1. Criação do BGP *speaker* com AS privado
2. Associação com a rede externa
3. Definição do *peer* (roteador *edge*) com AS público
4. Associação do BGP *speaker* com o *peer*
5. Associação do BGP *speaker* com um *network node*

NEUTRON-BGP-AGENT

1. neutron **bgp-speaker-create** --ip-version 6 --local-as 65001 --advertise-floating-ip-host-routes false **bgp1**
2. neutron **bgp-speaker-network-add** **bgp1** net-ext
3. neutron **bgp-peer-create** --peer-ip 2804:1234:1234::1 --remote-as 262545 **bgp1-peer1**
4. neutron **bgp-speaker-peer-add** **bgp1** **bgp1-peer1**
5. neutron **bgp-dragent-speaker-add** `openstack network agent list --agent-type bgp --host cldvcptr001.cloudmandic.com.br -c ID -f value` **bgp1**

CRIAÇÃO PILHA IPv6 – ADMIN



CRIAÇÃO PILHA IPv6 – ADMIN

```
openstack address scope create --share --ip-version 6 address-scope-ip6
```

```
openstack subnet pool create --address-scope address-scope-ip6 --pool-prefix  
2804:1234:1:1234::/64 --default-prefix-length 64 --max-prefix-length 64 ext-pool-ip6
```

```
openstack subnet create --network net-ext --ip-version 6 --subnet-pool ext-pool-ip6 -  
-no-dhcp --dns-nameserver 2001:4860:4860::8888 --allocation-pool  
start=2804:1234:1:1234::100,end=2804:1234:1:1234:ffff:ffff:ffff:ff00 sub6-ext
```

```
openstack subnet pool create --address-scope address-scope-ip6 --share --pool-  
prefix 2804:1234:1:1000::/52 --default-prefix-length 64 --max-prefix-length 64 --  
default default-pool-ip6
```

CRIAÇÃO PILHA IPv6 – PROJETO

openstack subnet create

```
--network net-int  
--ip-version 6  
--ip_ra_mode dhcipv6-stateless  
--ip_address_mode dhcipv6-stateless  
--subnet-pool default-pool-ip6  
--dns-nameserver 2001:4860:4860::8888  
sub6-int
```

IMPLEMENTAÇÃO

IPv6 NO OPENSTACK

IMPLEMENTAÇÃO PRÓPRIA

+ WEB Service:
Python Flask

+ BGP Speaker:
RYU SDN Controller

+ Uso de alarmes via aodh

+ Alta disponibilidade
2 agentes autônomos duplicados

AODH – DEFINIÇÃO DE ALARMES

/etc/ceilometer/events_definitions.yaml

```
- event_type: router.*  
  traits:  
    <<: *network_traits  
  resource_id:  
    fields: ['payload.router.id', 'payload.id']  
- event_type: router.interface.*  
  traits:  
    <<: *network_traits  
  port_id:  
    fields: payload.router_interface.port_id  
  subnet_id:  
    fields: payload.router_interface.subnet_id  
  router_id:  
    fields: payload.router_interface.id
```

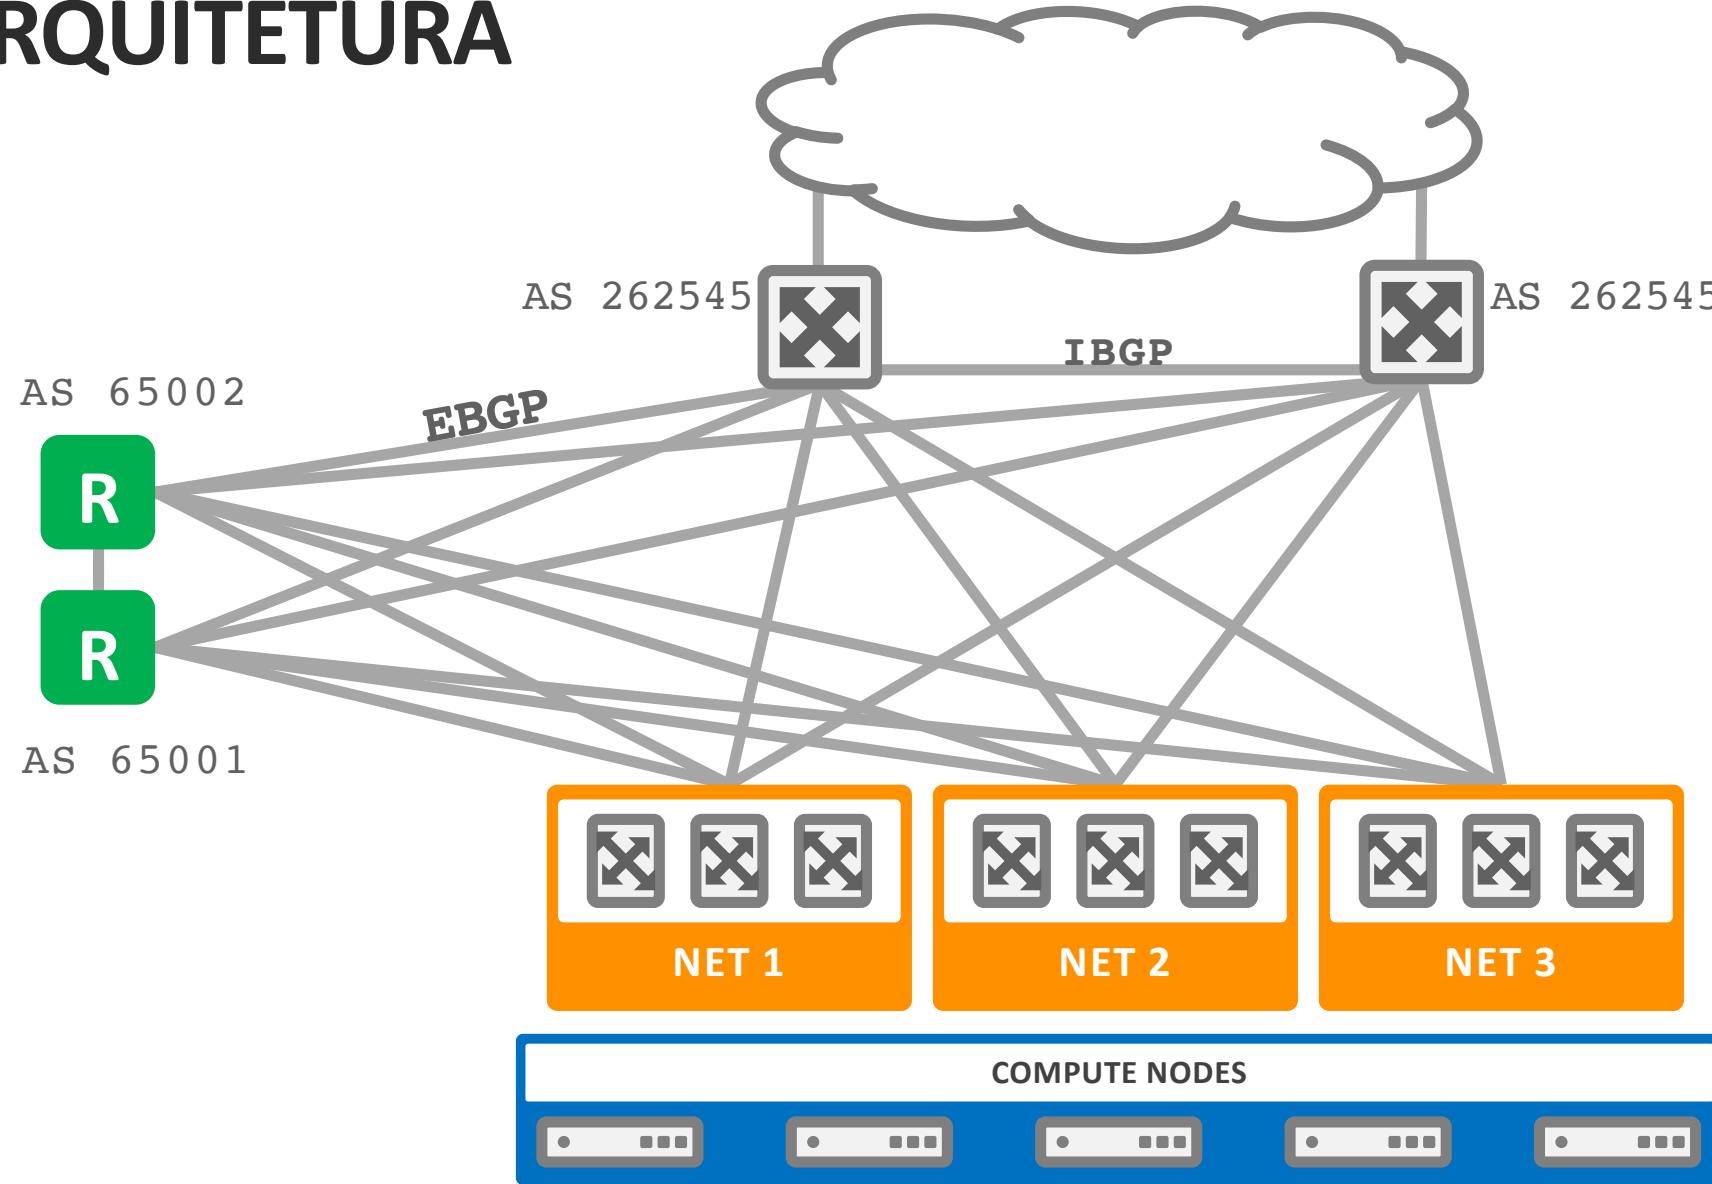
AODH – DEFINIÇÃO DE ALARMES

```
jose@jose-think:~$ openstack alarm show 991922bc-8b3b-435c-abb4-a74210c91889
+-----+-----+
| Field          | Value
+-----+-----+
| alarm_actions | [u'http://10.0.0.1:5125', u'http://10.0.0.2:5125']
| alarm_id       | 991922bc-8b3b-435c-abb4-a74210c91889
| description    | Alarm when router.interface.* event occurred.
| enabled        | True
| event_type     | router.interface.*
| insufficient_data_actions | []
| name           | CD123456-routing
| ok_actions     | []
| project_id     | 0b5930a580ea4cf9451a1ef2d6f71af
| query          |
| repeat_actions | True
| severity       | critical
| state          | insufficient data
| state_timestamp| 2018-04-06T21:43:45.750009
| time_constraints| []
| timestamp      | 2018-04-06T21:43:45.750009
| type           | event
| user_id        | 59c13f719bf24c11b66bacbb1efd839d
+-----+-----+
```

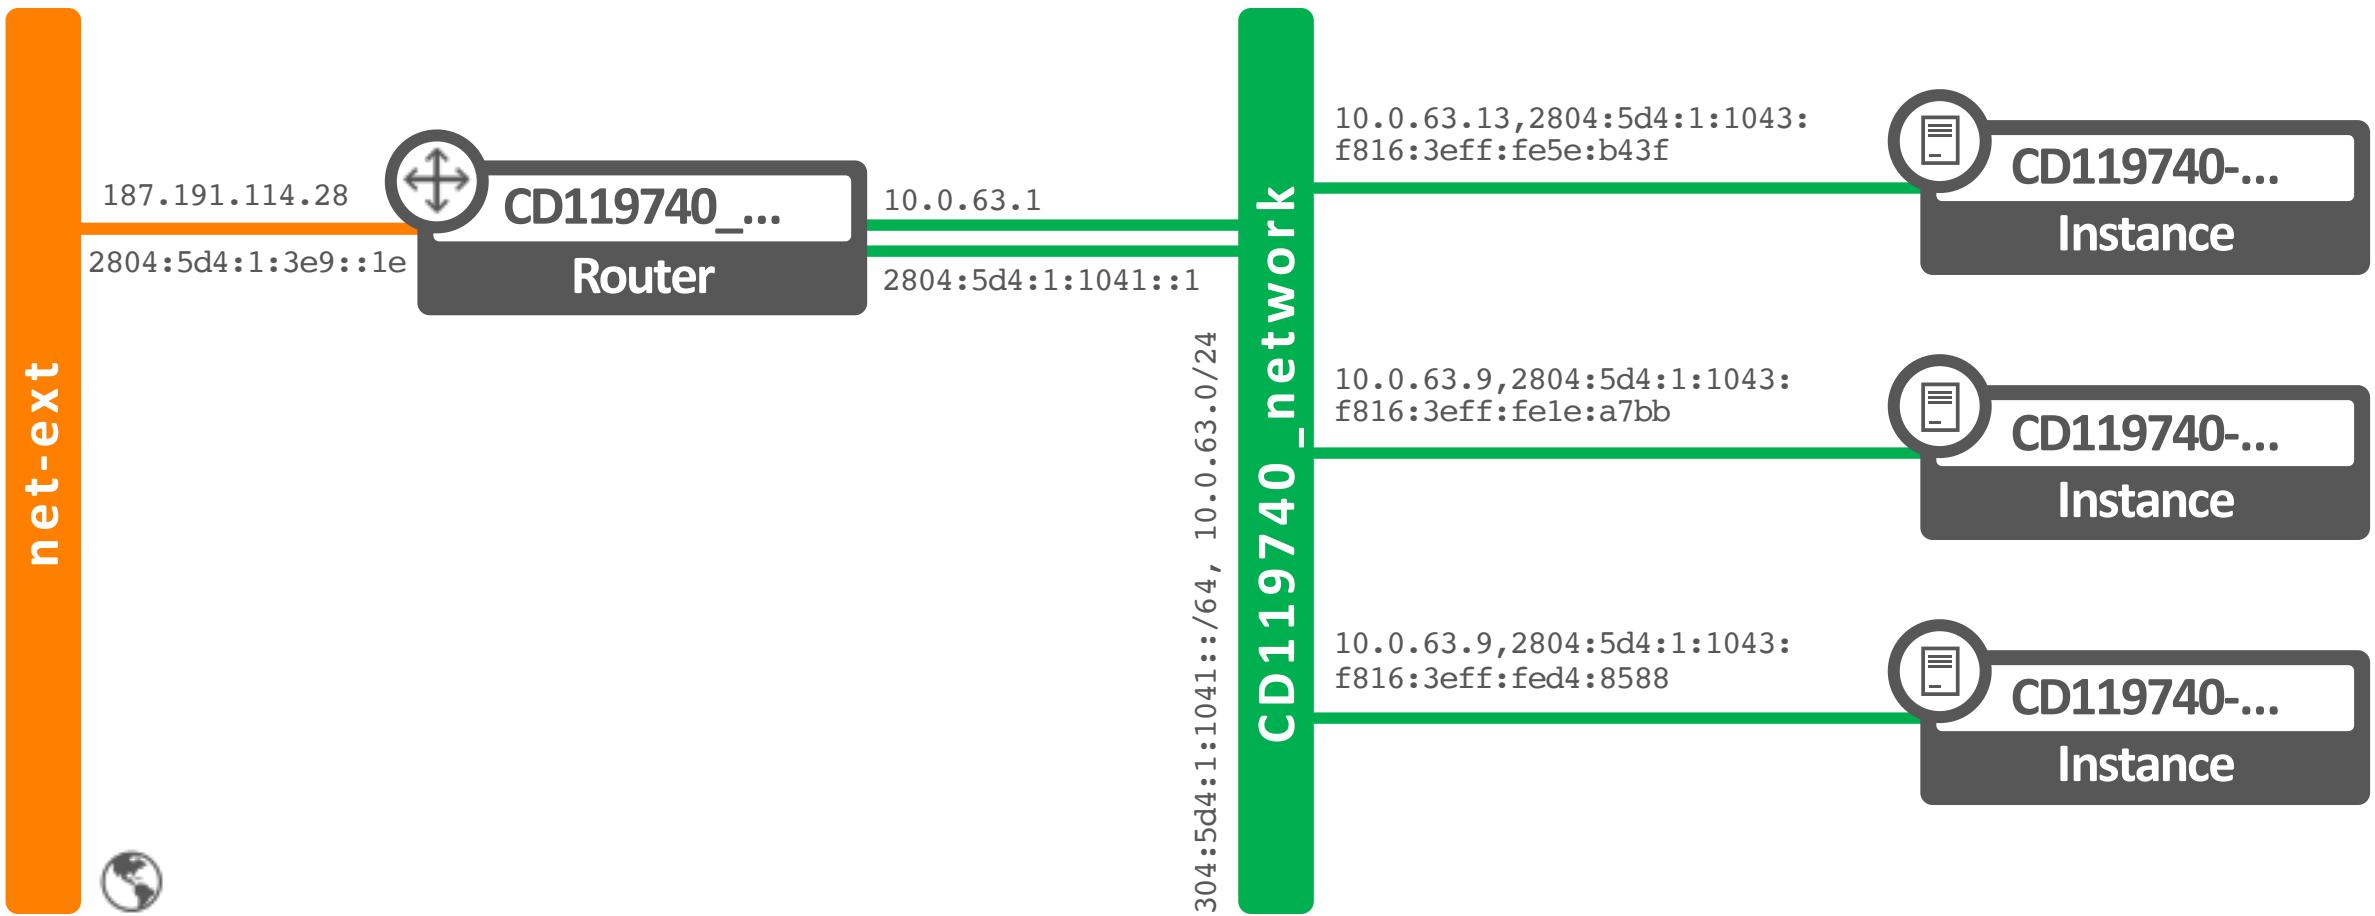
PROCESSO

1. O projeto é criado
2. É criado um alarme aodh para *router-add-subnet*
3. Um rede é criada para cada projeto
4. Um roteador virtual é criado para cada projeto
5. Duas subredes são criadas dentro da rede
 1. IPv4 a partir do *subnet pool IPv4*
 2. IPv6 a partir do *subnet pool IPv6*
6. O alarme notifica aos BGP *speakers*
7. O BGP speaker propaga a rota para os roteadores *edge*

ARQUITETURA



RESULTADO FINAL



RESULTADO FINAL

```
ubuntu@ub14d:~$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 scope host lo
            valid_lft forever preferred_lft forever
        inet6 ::1/128 scope host
            valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether fa:16:3e:96:e5:61 brd ff:ff:ff:ff:ff:ff
        inet 10.0.0.6/24 brd 10.0.0.255 scope global eth0
            valid_lft forever preferred_lft forever
        inet6 2804:6371:1:600:fa:6:3eff:fe96:e561/64 scope global dynamic
            valid_lft 86375sec preferred_lft 14375sec
        inet6 fe80::f816:3eff:fe96:e561/64 scope link
            valid_lft forever preferred_lft forever
```

CONSIDERAÇÕES: SECURITY GROUPS

Add Rule

Rule*

ALL ICMP

Direction

Ingress

Remote* ?

CIDR

CIDR ?
::/0

Manage Security Group Rules: default (a9947ba5-522e-4708-b662-d8c5af3c9544)

<input type="checkbox"/> Direction	<input type="checkbox"/> Ether Type	IP Protocol	Port Range	Remote IP Prefix
<input type="checkbox"/> Ingress	<input type="checkbox"/> IPv6	Any	Any	-
<input type="checkbox"/> Egress	<input type="checkbox"/> IPv6	Any	Any	::/0

LIMITAÇÕES

+ cloud-init

PRÓXIMOS PASSOS

- + DNSaaS, DNS6,
integração com **OpenStack Designate**
- + VTEP – Integração com **Cumulus**
- + VPNaas usando *containers*

CONCLUSÃO

- + Criação de subredes IPv6 no OpenStack
 - + Modos de atribuição de endereços
- + Conexão com a Internet
 - + IPv6 Prefix delegation
 - + Roteamento
- + Implementação
 - + *aodh*

LINKS

+ OpenStack IPv6 Networking

<https://docs.openstack.org/liberty/networking-guide/adv-config-ipv6.html>

+ BGP dynamic routing

<https://docs.openstack.org/newton/networking-guide/config-bgp-dynamic-routing.html>

+ Prefix delegation

<https://docs.openstack.org/newton/networking-guide/config-ipv6.html#prefix-delegation>

JOSÉ CASTILLO LEMA

DIRETORIA DE OPERAÇÕES E SOLUÇÕES AO CLIENTE

especialistascloud@mandic.net.br



Especialistas em Clouds.

Mandic Clouds:



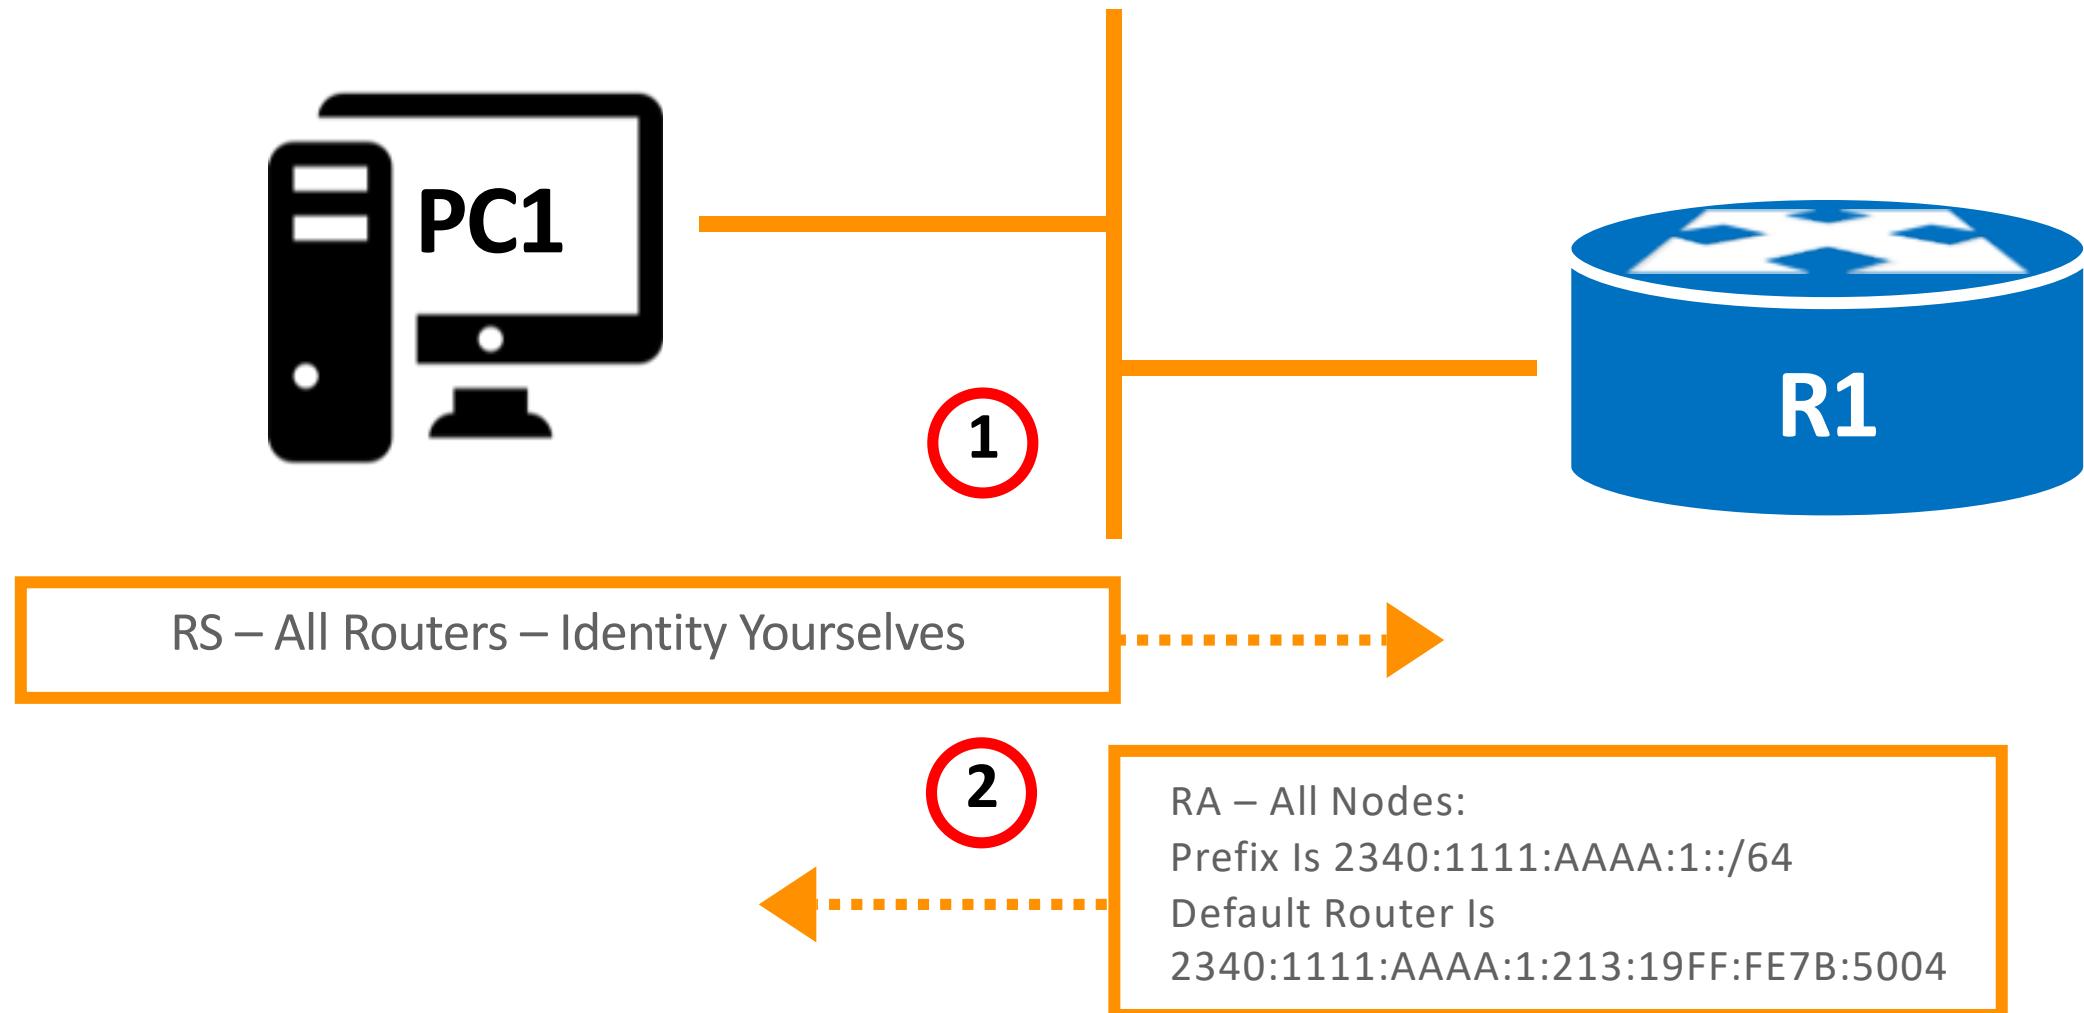
vmware®

openstack.

Azure



NEIGHBOR DISCOVERY PROTOCOL – NDP



OPENSTACK NEUTRON – PROPRIEDADES IPv6

ipv6_address_mode

- + determina como os endereços IPv6 são gerados e assignados
- + *slaac / dhcpv6-stateless / dhcpv6-stateful*

ipv6_ra_mode

- + determina como os roteadores virtuais enviam RAs
- + *slaac / dhcpv6-stateless / dhcpv6-stateful*

OPENSTACK NEUTRON – PROPRIEDADES IPv6

ipv6 ra mode	ipv6 address mode	radvd A,M,O	External Router A,M,O	Description
slaac	slaac	1,0,0	Off	Guest instance obtains IPV6 address from OpenStack managed radvd using SLAAC.
dhcpv6-stateful	dhcpv6-stateful	0,1,1	Off	Guest instance obtains IPV6 address from dnsmasq using DHCPv6 stateful and optional info from dnsmasq using DHCPv6
dhcpv6-stateless	dhcpv6-stateless	1,0,1	Off	Guest instance obtains IPv6 address from OpenStack managed radvd using SLAAC and optional info from dnsmasq using DHCPv6
Texto...	Texto...	Texto...	Texto...	Texto...