

DEVCONF.cz

Shifting Performance Engineering Left

Building Performance into the Development Cycle



kube-burner

Raúl Sevilla
rsevilla@redhat.com

Jose Castilo Lema
jlema@redhat.com

whoarewe

- 8 years in Red Hat
- 5 years in the OCP Perf&Scale Team
- 3 years in the Iberia Services Team



Raúl
Sevilla



JOSE
CASTILLO LEMA

- 6 years in Red Hat
- 4 years in perf/scale department
- 2 years as Telco Cloud Architect in the Solutions & Technology Practices team

What we'll discuss today

1. The problem
2. Shift-left engineering
3. Continuous Performance Testing (CPT)
4. Workflow
5. Tooling
 - a. Kube-burner
 - i. Demo
6. Case studies
7. Challenges and next steps

Continuous Performance Testing

The problem

- ▶ Performance testing often occurs **too late in the pipeline**
- ▶ Late detection of performance issues increases **cost**
- ▶ Real-world consequences: downtime, latency, etc.


Performance testing – the past

OpenShift product lifecycle



At best we normally would fit in at the end of the development cycle

Shift-left engineering



Development



Shift left

**Performance
Engineering**

- Testing
- Tuning
- Capacity planning



Shift right



Operations

- Monitoring
- Capacity planning

Alexander Podelko, *Continuous Performance Testing: Challenges and approaches*


What is CPT?

Continuous Performance Testing (CPT) is the practice of executing performance benchmarks (tests) continuously throughout the development lifecycle.

The execution of said Performance benchmarks is **orchestrated in an automated pipeline** (like Jenkins, Prow, etc.), preferably in the pipeline your developers are working out of.

CPT goals

OpenShift product lifecycle



Where we should fit in the development cycle

Why CPT?

- Shifts Performance testing left
- Earlier feedback enables faster fixes
- Reduces overall testing costs
- Improves developer ownership of performance
 - Continuous Improvement Mindset
- More time for technical improvements
- Allows us to drastically improve the coverage matrix and frequency of the testing
 - More data → Need for automatic **performance regression frameworks**
 - **Change point detection** mechanisms

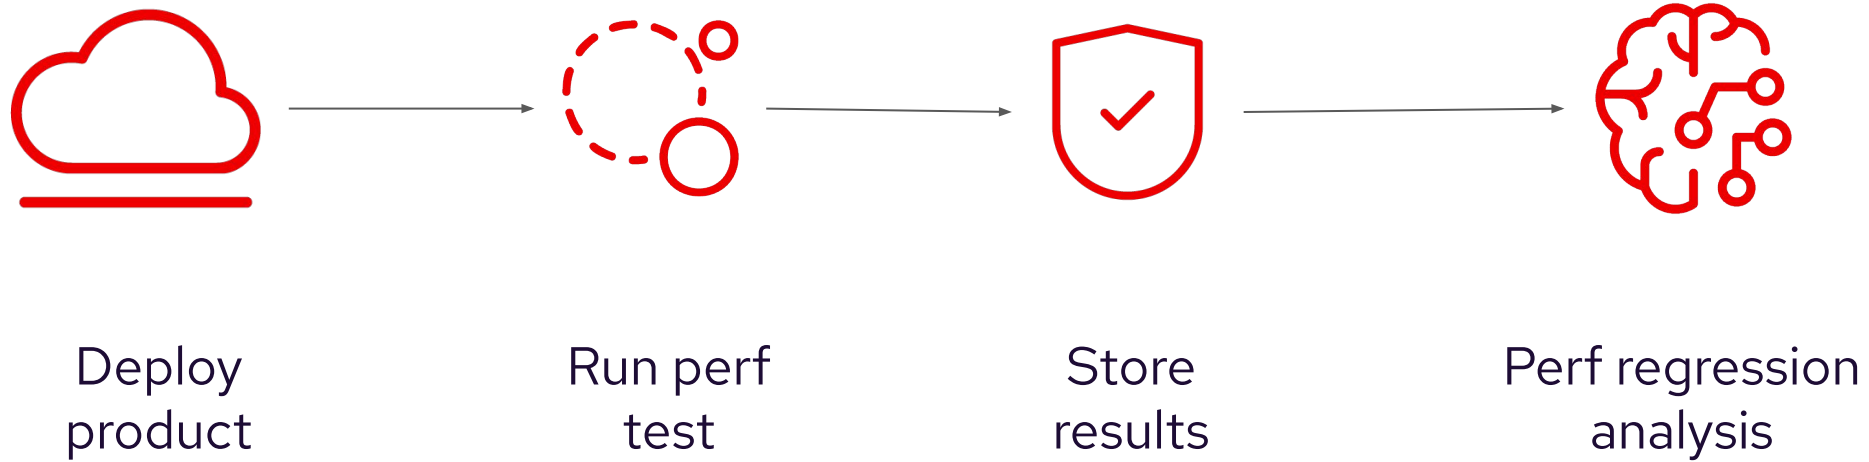
CPT – Cultural shift

- From siloed to **dev + Performance collaboration**
- Embed performance mindset in development sprints
- Performance as part of 'Definition of Done'

Key principles to **navigate change**:

- Document workflows and KPI's
- Collaboration with QE counterparts is key
 - Handoffs
 - Responsibility matrix

CPT workflow



Achievements

- Catch performance/scalability regressions in an early stage
- Redefined a new test/platform coverage matrix standard

5 platforms

AWS, Azure, GCP, IBM, Baremetal +
Managed services offerings: ROSA and ARO

+100 weekly tests

Covering nightly, early candidates, release candidates,
stable releases and long term support

Platform +

- OpenShift Virtualization
- Layered products (OLS, ACS, Kepler, ...)

Tooling - Kube-burner

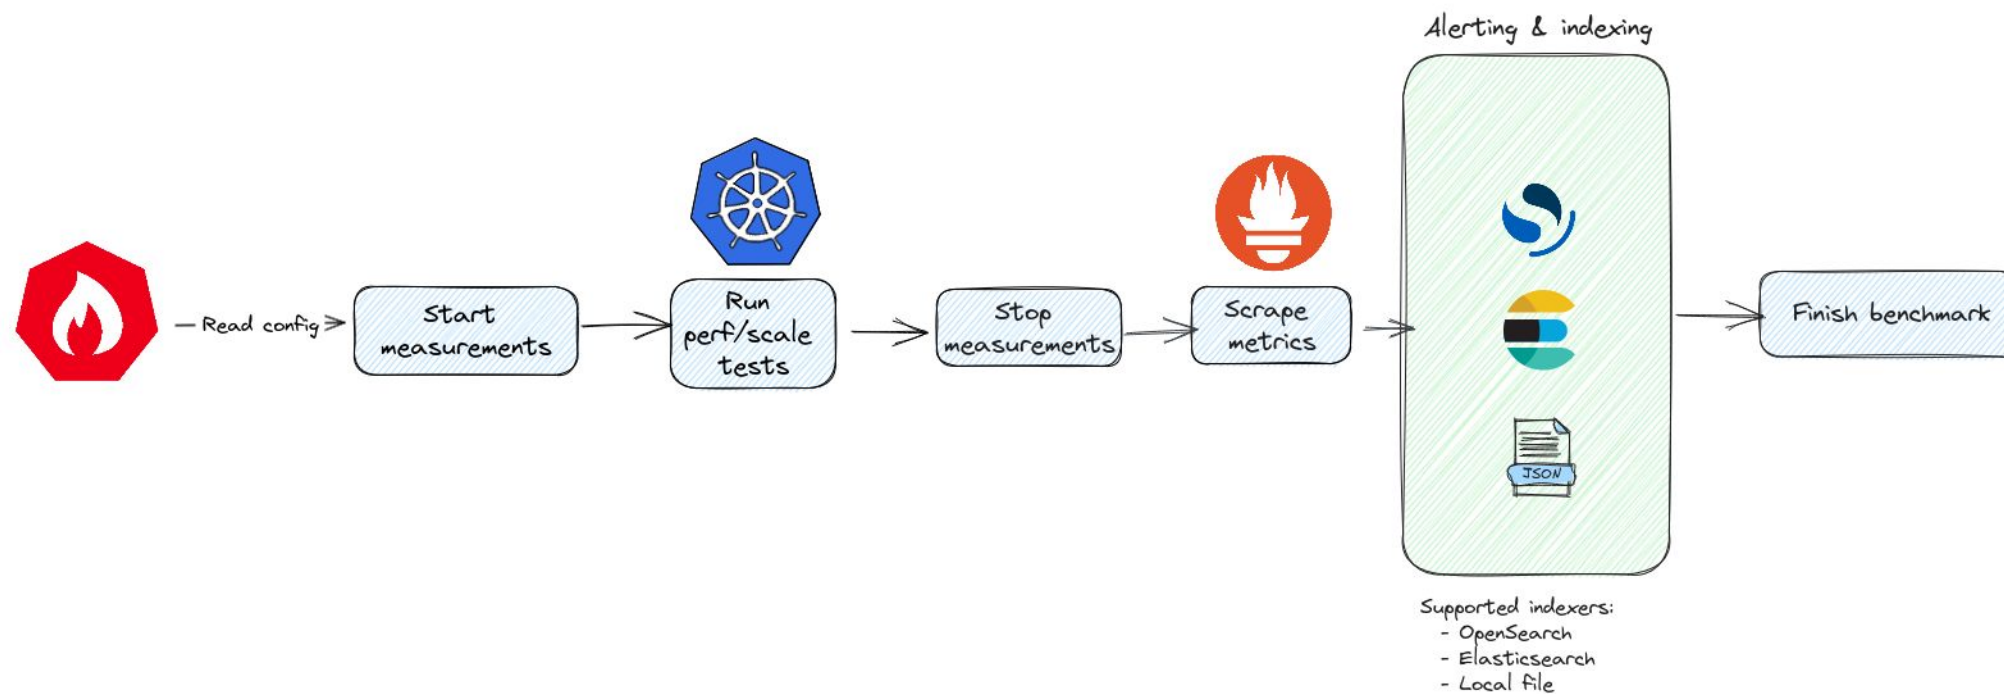
What's this?

Kubernetes **performance and scale test orchestration framework**.

It provides multi-faceted functionality, some most important of which are:

- CRUD operations over Kubernetes resources at scale
- Prometheus metric collection and indexing
- Measurements
- PromQL based alerting

Kube-burner



Kube-burner Highlights

- Configurable API call rate through **QPS/Burst**
- Easy to use, friendly **YAML** based configuration files
- Very flexible configuration, with lots of options and customizations
- Fully compatible with vanilla K8s and other K8s distros
- **Canned workloads** and more customizations can be added through plugins
- **CNCF**'s Sandbox project and currently working on upgrading it to Incubating

Kube-burner jobs

Support for different **job types**

- **Create:** Creates the listed objects
- **Delete:** Deletes objects enumerated in the objects list by using label selectors
- **Patch:** Patches objects with a predefined patch template
- **Read:** Lists objects described in the objects list.
- **Kubevirt:** Can be used to execute virtctl type commands (using kubectl) over running VMs.

Kube-burner observability

- **Prometheus** metrics scraping
 - Collection of Prometheus metrics over the workload runtime
 - Evaluates Prometheus expressions to fire alerts
- **Measurements**
 - Additional observability w/o using Prometheus:
 - Pod/Job/PVC Latency
 - VM/VMI latency (Kubevirt)
 - Pprof
 - And more

Kube-burner alerting

Evaluates Prometheus expressions configured in a file called **alert-profile**.

When the PromQL condition is met kube-burner fires an alert with a defined severity:

- *info, warning, error or critical*

These alerts can be indexed along with the collected metrics or measurement.

In the workloads currently executed within CPT, just a few alerts have an error severity or higher

Kube-burner-ocp

Some of the [supported workloads](#) are:

- **Cluster-density-v2:** Creates a high number of objects that generate a significant load in the control plane and the CNI plugin. (Deployments, secrets, builds, routes, probes, etc)
- **Node-density:** Fills with pause pods all the worker nodes of the cluster
- **Node-density-cni:** Fills with client/server (curl and nginx) pods sending periodic traffic through services
- **Node-density-heavy:** Fills with client/server (server is a postgresql database), where the client performs periodic queries in the database
- **Virt-density:** Spawns a series of tiny VMs across the worker nodes of the cluster
- **Virt-density-udn:** Similar to the previous, but those VMs are configured to use OVN's UDN interfaces.
- **Virt-migration:** Tests how the cluster and the storage backend handles mass migration of VMs

It's all about data – Automatic regression detection

When a benchmark finishes, it doesn't mean a success from a Perf&Scale perspective.

We post-process the benchmark's generated data using our own **regression detection framework**, [Orion](#).

- Uses different algorithms to detect regressions and deviations in some specific metrics we consider reliable and stable.

We've been able to early detect, report and fix (when possible) many performance regressions. i.e:

- A 25% deviation in the pod latency can mean many things, since lots of components are involved in a pod creation
 - We narrow down the issue by snooping around the updated bits in that specific OpenShift build.

Success stories – KAPI latency

time	uuid	buildUrl	readOnlyAPICallsLatency_max
2025-01-13 11:38:02 +0000	a485ef5b-..-e6d6b35532ca	https://prow...aws-4.19-..payload-control-plane-6nodes/1878744929488867328	0.984
2025-01-14 19:20:40 +0000	95446d27-..-f2d8d7d8ce27	https://prow...aws-4.19-..payload-control-plane-6nodes/1879221652647055360	0.984
2025-01-15 03:28:01 +0000	cc0b1a6d-..-68a59992b6b8	https://prow...aws-4.19-..payload-control-plane-6nodes/1879343743165796352	0.984
2025-01-15 07:43:23 +0000	20058cc2-..-7009402a6c13	https://prow...aws-4.19-..payload-control-plane-6nodes/1879411213251645440	0.976
2025-01-16 01:26:21 +0000	cdf21278-..-71782e3ab17e	https://prow...aws-4.19-..payload-control-plane-6nodes/1879675626953117696	0.984
2025-01-16 08:26:51 +0000	0b4d30e2-..-e056c094a622	https://prow...aws-4.19-..payload-control-plane-6nodes/1879784780980031488	0.984
2025-01-16 16:23:04 +0000	bfb0efa0-..-d1b0cda6dd41	https://prow...aws-4.19-..payload-control-plane-6nodes/1879902896481374208	0.9
2025-01-17 00:08:31 +0000	2d6b598f-..-64d5c98e4236	https://prow...aws-4.19-..payload-control-plane-6nodes/1880020138195947520	3.98
		
			+923.1% !!
		
2025-01-17 06:37:24 +0000	391817f4-..-557cd13bd32e	https://prow...aws-4.19-..payload-control-plane-6nodes/1880118891967942656	19.9
2025-01-17 11:39:51 +0000	db926aaf-..-da59c974574a	https://prow...aws-4.19-..payload-control-plane-6nodes/1880193587291885568	19.9
2025-01-17 18:23:02 +0000	5948ce33-..-94cc8cb517fb	https://prow...aws-4.19-..payload-control-plane-6nodes/1880291350289584128	3.94
2025-01-18 00:51:31 +0000	efa220e8-..-60e4876a3e52	https://prow...aws-4.19-..payload-control-plane-6nodes/1880388538206261248	19.6
2025-01-18 06:12:03 +0000	57aff787-..-28e24cf051f7	https://prow...aws-4.19-..payload-control-plane-6nodes/1880474335618011136	28.6
2025-01-18 12:57:00 +0000	11e53cf9-..-9f77ebfc9bdc	https://prow...aws-4.19-..payload-control-plane-6nodes/1880574934015545344	3.86
2025-01-18 18:16:42 +0000	82069eee-..-b3d06c2a1fa8	https://prow...aws-4.19-..payload-control-plane-6nodes/1880655151874707456	3.74

Success stories – KAPI latency

- Focus on one single build containing **618 commits** (!)
- Kernel team took bisection down through all commits to land on a single (or two) change(s)


The **madvise()** system call is used to give advice or directions to the kernel about the address range beginning at address *addr* and with size *size*. **madvise()** only operates on whole pages, therefore *addr* must be page-aligned. The value of *size* is rounded up to a multiple of page size. In most cases, the goal of such advice is to improve system or application performance.

MADV_RANDOM

Expect page references in random order. (Hence, read ahead may be less useful than normally.)

Success stories – Pod Latency

- Focus on [a specific Pull Request](#) containing **2839 commits** (!)
- Git **bisect**, recompile kubelet and run test to find out the specific "**bad commit**"



Kube-burner-ocp demo

More **demos** at <https://github.com/kube-burner/kube-burner-demos>

Challenges

- Resistance to change and lack of expertise
- Tooling integration issues
- We can only detect regressions on metrics we are tracking
- Avoid **duplicated testing**
 - Reduce redundant tested combinations
- Enable developers to **interpret** (and predict) performance results
 - Exploring the use of AI to help* analyze results

Solutions

Key principles to **navigate change**:

- Training
- Automation
- Establish a solid protocol for dealing with test failures and flakiness
- Champion roles

Roadmap – where to start

- **Start with one service**

- Educate teams on performance fundamentals
- Build integrated performance dashboards into CI/CD pipelines
- Include performance in code reviews

Food for thought:

- How can performance testing fit into your current sprint workflow?
- What's one thing you could shift left today?

Wrap up

1. Shift-left engineering
2. Continuous Performance Testing (CPT)
3. Workflow
4. Tooling
 - a. kube-burner
5. Real life case study
 - a. OCP → CoreOS → Kernel → syscall
6. Challenges and next steps

Q & A

