

# AI-Powered Performance Insights

Integrating OVS/OVN automatic performance regression analysis with LLMs

**Joe Talerico, Mohit Sheth, Raul Sevilla, José Castillo**  
Perf&Scale Engineering

# whoarewe



**JOSE  
CASTILLO LEMA**

- 6 years in Red Hat
- 4 years in perf/scale department
- 2 years as Telco Cloud Architect in the Solutions & Technology Practices team

- 10 years in Red Hat
- 6 years in the OCP Perf&Scale Team
- 4 years in the Iberia Services Team



**Joe  
Talerico**



**Raúl  
Sevilla**



# Agenda

1. The problem
2. Shift-left engineering
3. Continuous Performance Testing (CPT)
4. Workflow
5. Tooling
6. How does AI fit in?
7. Challenges and next steps

## The problem

- ▶ Performance testing often occurs **too late in the pipeline**
- ▶ Late detection of performance issues increases **cost**
- ▶ Real-world consequences: downtime, latency, etc.

# Performance Testing – The past


Red Hat Confidential, I guess?

## OpenShift product lifecycle




At best we normally would fit in at the end of the development cycle

# Shift-left engineering




**Development**



**Shift left**

## Performance Engineering

- Testing
- Tuning
- Capacity planning



**Shift right**



**Operations**


- Monitoring
- Capacity planning

## What is CPT?

**Continuous Performance Testing** (CPT) is the practice of executing performance benchmarks (tests) continuously throughout the development lifecycle.

The execution of said Performance benchmarks is **orchestrated in an automated pipeline** (like Jenkins, Prow, etc.), preferably in the pipeline your developers are working out of.

# Workflow




# Why CPT?

- ▶ Shifts Performance testing left
- ▶ Earlier feedback enables faster fixes
- ▶ Reduces overall testing costs
- ▶ Improves developer ownership of performance → Continuous Improvement

## Mindset

- ▶ More time for technical improvements
- ▶ Allows us to drastically improve the coverage matrix and frequency of the testing
  - More data → Need for automatic **performance regression frameworks**
    - **Change point detection** mechanisms

## OpenShift product lifecycle



After adopting CPT, where we fit into the lifecycle.  
We would like to see how we can enable the same process upstream!

## Cultural shift

- ▶ From siloed to **dev + Performance collaboration**
- ▶ Embed performance mindset in development sprints
- ▶ Performance as part of 'Definition of Done'

Key principles to **navigate change**:

- ▶ Document workflows and KPI's
- ▶ Collaboration with QE counterparts is key
  - Handoffs
  - Responsibility matrix

# PerfScale x Openshift Networking

## 2025 in review

- ▶ Net New Workloads added (Control & Data-path)
  - UDN L2 & L3
  - Virt UDN
  - Networkpolicy
  - Egress IP
  - BGP
- ▶ 30+ downstream bugs opened
- ▶ Improved tooling to identify regressions in a PR
- ▶ BiWeekly engagement with developers

## Achievements

- Catch performance/scalability regressions in an early stage
- Redefined a new test/platform coverage matrix standard

5 platforms

AWS, Azure, GCP, IBM, Baremetal +  
Managed services offerings: ROSA and ARO

+100 weekly tests


Covering nightly, early candidates, release candidates,  
stable releases and long term support

Platform +

- OpenShift Virtualization
- Layered products (OLS, ACS, Kepler, ...)


# Control-Plane Tooling

Red Hat Confidential - Associates only



# Data-Plane Tooling - k8s-netperf

Red Hat Confidential - Associates only



## k8s-netperf:


- Deploys benchmark assets in the given cluster
- Drives benchmark orchestration and results collection
- Normalizes and index benchmark results along with Prometheus metrics in local files or ElasticSearch/OpenSearch
- Garbage collects benchmark assets

Supports throughput and latency benchmarks:

- Pod 2 pod via SDN or hostNetwork (including VMs)
- Pod 2 service
- Multiple drivers: netperf, uperf, iperf3, ib\_write\_bw (RoCE traffic)

# Data-Plane Tooling - ingress-perf

Red Hat Confidential - Associates only




## Ingress-perf:

- Deploys benchmark assets in the given cluster
- Drives benchmark orchestration and results collection
- Normalizes and index benchmark results along with Prometheus metrics in local files or ElasticSearch/OpenSearch
- Garbage collects benchmark assets

# Regression Analysis – Tooling

Red Hat Confidential - Associates only

Across all of our Performance and Scale Tooling, there is a common workflow.



With this common workflow and data warehouse, we can use a common tool for the regression analysis, called [Orion](#)

Orion works by querying the data warehouse for all results that *match a specific fingerprint*. Once we have all the matches, Orion will then collect the desired system metrics the user provided and run the provided regression analysis algorithm against the dataset – We use Apache Otava (e-divisive means algorithm, [learn more](#))

<sup>17</sup>

# Regression Analysis – Tooling

Red Hat Confidential - Associates only

Example of an Orion configuration:

```
tests :  
  - name : 24-node-scale-cdv2  
    index: {{ es_metadata_index }}  
    benchmarkIndex: {{ es_benchmark_index }}  
    metadata:  
      platform: AWS  
      clusterType: self-managed  
      masterNodesType: m6a.xlarge  
      masterNodesCount: 3  
      workerNodesType: m6a.xlarge  
      workerNodesCount: 24  
      benchmark.keyword: cluster-density-v2  
      ocpVersion: {{ version }}  
      networkType: OVNKubernetes  
      jobType: {{ jobtype | default('periodic') }}  
    not:  
      stream: okd
```

**This example fingerprint will retrieve:**

- all runs on AWS
- m6a.xlarge control-plane
- m6a.xlarge workers
- 24 workers
- The benchmark was cluster-density-v2

# Regression Analysis - Tooling

Red Hat Confidential - Associates only

Example of an Orion OVN analysis:

ocpVersion	podReadyLatency_P99	ovsCPU-Workers_avg	ovsCPU-irate-all_avg
4.20.0-0.nightly-2025-11-04-042419	4000	2904.67	0.23164
4.20.0-0.nightly-2025-11-05-022131	5000	2902.71	0.228098
4.20.0-0.nightly-2025-11-05-192504	4000	2960.68	0.235272
4.20.0-0.nightly-2025-11-06-033058	4000	2907.6	0.239161
4.20.0-0.nightly-2025-11-06-091216	4000	2934.81	0.236172
4.20.0-0.nightly-2025-11-06-175730	4000	2973.54	0.239612
4.20.0-0.nightly-2025-11-07-105814	4000	2775.94	0.246011
4.20.0-0.nightly-2025-11-11-060054	5000	2956.15	0.244239
4.20.0-0.nightly-2025-11-11-124009	4000	2926.77	0.237181
4.20.0-0.nightly-2025-11-11-205214	4500	2661.63	0.261356
4.20.0-0.nightly-2025-11-12-034103	4000	2646.27	0.23104
4.20.0-0.nightly-2025-11-12-111016	4000	2636.41	0.234695
4.20.0-0.nightly-2025-11-12-185351	4000	2661.88	0.230927
4.20.0-0.nightly-2025-11-13-021810	4000	2637.28	0.244027
4.20.0-0.nightly-2025-11-13-101159	4000	2603.91	0.234688
4.20.0-0.nightly-2025-11-13-165930	4000	2663.08	0.242627
4.20.0-0.nightly-2025-11-13-205930	5000	2657.42	0.258178
4.20.0-0.nightly-2025-11-14-043920	4000	2623.96	0.245345
4.20.0-0.nightly-2025-11-14-130157	4000	2628.41	0.25806
4.20.0-0.nightly-2025-11-14-213105	4000	2625.4	0.239047
4.20.0-0.nightly-2025-11-15-035645	4000	2658.08	0.234444
4.20.0-0.nightly-2025-11-15-103436	4000	2662.73	0.255973
4.20.0-0.nightly-2025-11-15-175712	4000	2666.94	0.237714
4.20.0-0.nightly-2025-11-16-023241	4000	2637.71	0.2352
4.20.0-0.nightly-2025-11-16-090846	4000	2623.74	0.218506
4.20.0-0.nightly-2025-11-16-144116	4000	2649.8	0.242681
4.20.0-0.nightly-2025-11-16-202036	4000	2637.66	0.231949
4.20.0-0.nightly-2025-11-17-050836	4000	2591.05	0.244409
4.20.0-0.nightly-2025-11-17-135315	4000	2684.78	0.24645

\*Cropped due to size


# Where is the AI?

Red Hat Confidential - Associates only

Now that we have established our workloads and methodologies... The obvious question is, where is the AI?


We have recently introduced [Orion-MCP](#), which allows developers and users to ask their favorite LLM that is connected to our MCP server, "Has OpenShift Regressed?"

This isn't limited to simple requests, users can ask more complex questions, like "Has my PR introduced a regression" or for the LLM if OpenShift has regressed to look through the PR changelog and determine what PR could of possibly introduced the regression.



# Orion-mcp with LLM - Demo

Red Hat Confidential - Associates only



# Orion-mcp slack integration

Red Hat Confidential - Associates only



**PerfScale Padawan** APP 13 minutes ago

Analyzing PR performance... This may take a few moments.

PR Performance Analysis (AI generated)

## Performance Impact Assessment

- Overall Impact: Improvement
- Significant improvements ( $\geq 10\%$ ):
  - multusCPU\_avg (cluster-density)
- Moderate regressions (5-10%):
  - etcdCPU\_avg (crd-scale)
  - podReadyLatency\_P99 (node-density)
  - ovsCPU-Workers\_avg (node-density-cni)
  - apiserverCPU\_avg (crd-scale)
- Moderate improvements (5-10%):
  - ovnCPU-ovncontroller\_avg (node-density-cni)
  - multusCPU\_avg (node-density-cni)
  - ovnCPU-northd\_avg (node-density-cni)
  - ovsCPU-irate-all\_avg (cluster-density)
  - ovnMem-sbdb\_avg (node-density-cni)

## Most Impacted Metrics

Config: cluster-density

Metric	Baseline	PR Value	Change (%)
multusCPU_avg	0.14	0.13	-10.98
ovsCPU-irate-all_avg	0.15	0.15	-5.29
etcdCPU_avg	3.69	3.80	3.00
apiserverCPU_avg	5.06	4.98	-1.70
ovsMemory-Masters_max	167.07M	164.58M	-1.49
monitoringCPU_avg	1.03	1.05	1.38
ovnCPU_avg	1.48	1.47	-1.19
kubelet_avg	24.88	25.12	0.96
ovsMemory-Workers_max	469.41M	466.72M	-0.57
ovsMemory-all_avg	63.34M	63.06M	-0.46

Config: node-density-cni

Metric	Baseline	PR Value	Change (%)
ovnCPU-ovncontroller_avg	3.02	2.74	-9.43
multusCPU_avg	0.91	0.84	-7.80
ovnCPU-northd_avg	1.87	1.73	-7.50
ovsCPU-Workers_avg	2816.36	2984.33	5.96
ovnMem-sbdb_avg	18.19M	17.23M	-5.27
ovnkCPU-overall_avg	4.04	3.87	-4.13
ovnCPU-sbdb_avg	0.79	0.76	-4.11
ovnCPU-ovnk-controller_avg	26.14	25.20	-3.58
ovnCPU-nbdb_avg	0.65	0.63	-2.80
etcdCPU_avg	6.10	6.27	2.76

## Thread

X

Red Hat Confidential - Associates only



**jtaleric** 1 minute ago

OVSOVN Conf - Slackbot Thread (edited)

B I U S | ⌂ ⌃ ⌄ ⌅ ⌆ ⌇ ⌈ ⌉ ⌊ ⌋

Reply...

Also send to **bugzooka-test**

+ Aa 😊 @ | 📁 💬 | ✎



▼

X

# Orion-mcp metadata (wip)

Red Hat Confidential - Associates only

> Show me the **ovsCPU workers usage** of the last 5 results of OCP 4.19 on the daily virtualization bare metal runs and display the **corresponding OVN, OVN-k8s and OVS versions**, without images.

Run	ovsCPU workers	OVN	OVN-k8s	OVS	OCP
1 (most recent)	2904	<b>25.09</b>	<b>1.1</b>	<b>3.11</b>	4.21
2	2910	<b>25.09</b>	<b>1.1</b>	<b>3.11</b>	4.21
3	2897	<b>25.03</b>	<b>1.0</b>	<b>3.11</b>	4.20
4	2910	<b>24.11</b>	<b>1.0</b>	<b>3.9</b>	4.19
5	2900	<b>24.11</b>	<b>1.0</b>	<b>3.9</b>	4.19

# Challenges and next steps

Red Hat Confidential - Associates only

- ▶ Consider adding to the ovn-ci tests one of our performance workloads to start identifying performance and scale issues **upstream** in ovn-kubernetes?
- ▶ **Agentic** AI workload for CPT
- ▶ Increase **test coverage**
  - I.e.: localnets, etc.
- ▶ More **metadata** context
- ▶ More/better **integrations**:
  - Slack
  - Github

# Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



[linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)



[youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)



[facebook.com/redhatinc](https://www.facebook.com/redhatinc)



[twitter.com/RedHat](https://twitter.com/RedHat)